

# Reinforcement Learning for Particle Accelerators

## An Introduction

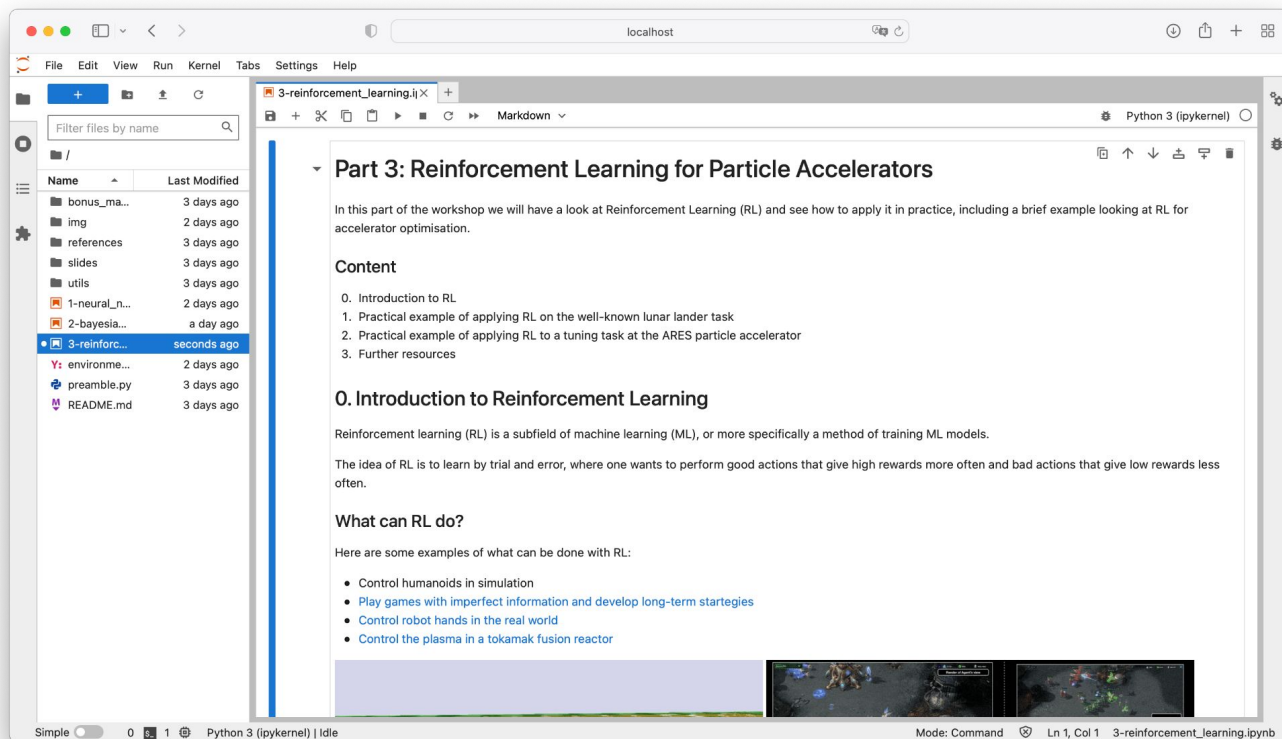
**Jan Kaiser and Oliver Stein**

*MT-ARD-ST3 pre-meeting ML workshop*



# Try Reinforcement Learning Yourself

Jupyter Notebook with code for examples from this presentation




# What can RL do?

## Some examples


Control humanoid in simulation




Control robot hands in the real world




Play games with imperfect information and develop long-term strategies



Control the plasma in a tokamak fusion reactor




View from inside the tokamak



Plasma state reconstruction

# Concepts of Reinforcement Learning

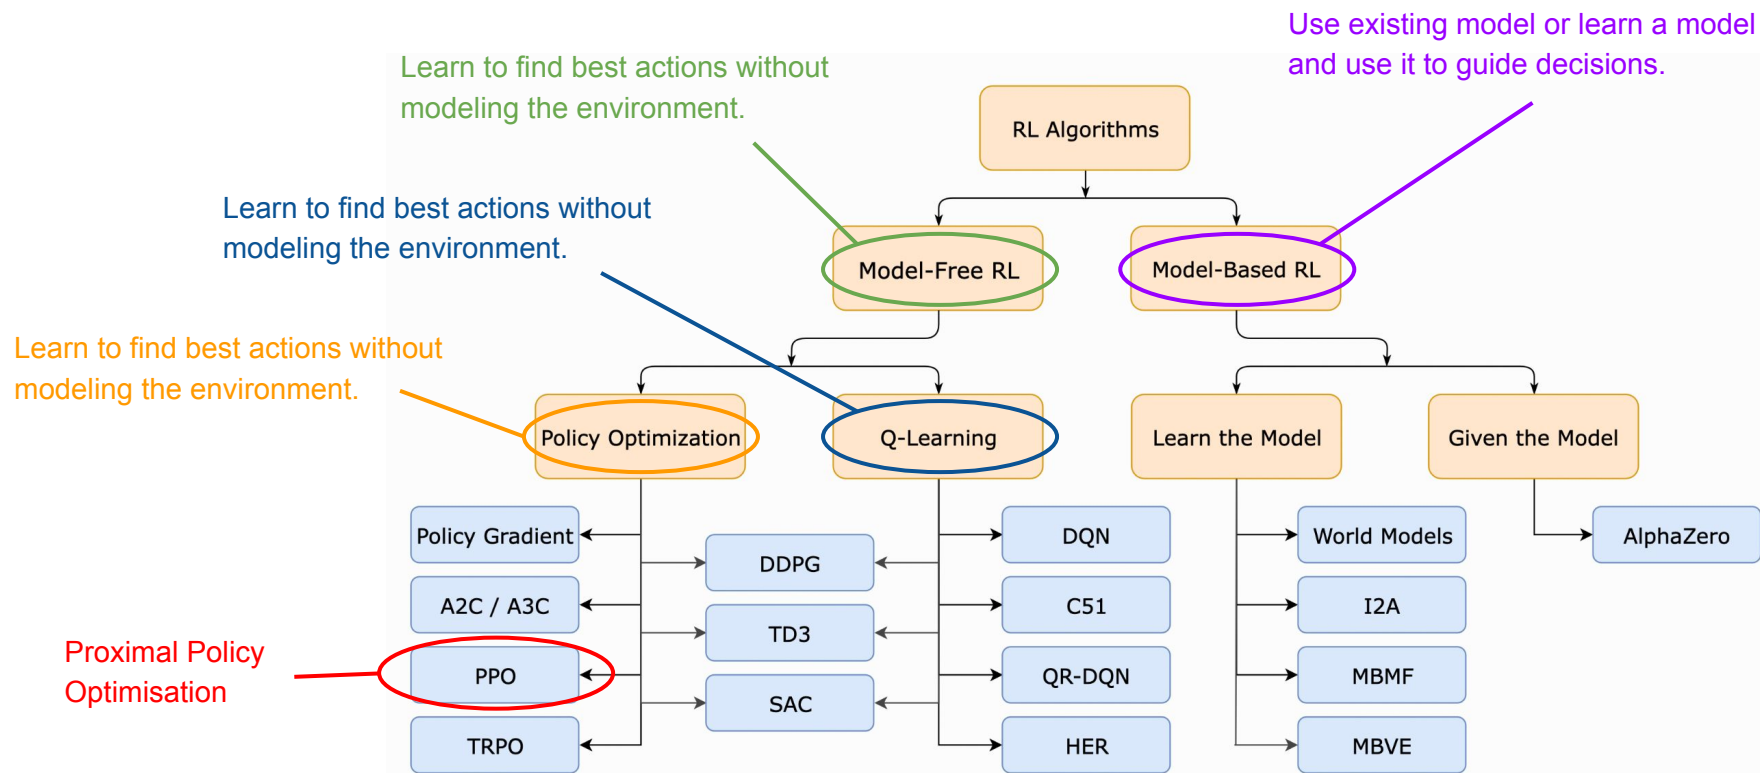
## Some examples



- The **agent** (or **policy**) is the function we are trying to learn and tells us what to do to solve the task.
- The **environment** is the world that the RL agent lives in and defines the task.
- **Actions** are how the RL agent interacts with the environment.
- **Observations** are what the agent sees of the environment.
- The **reward** is returned by the environment after each action and describes the goodness of that action.
- The **return** is the cumulative reward over time. The goal of RL is to maximise the return.


# Taxonomy of Reinforcement Learning

## A brief overview



# Lunar Lander Example


## Introduction




# Lunar Lander Example

## Actions and observations

Action




Observation




# Lunar Lander Example


## Rewards




-0.3




-0.03




+/- d



-100



+200







+10 each




# Lunar Lander Example

## Training results



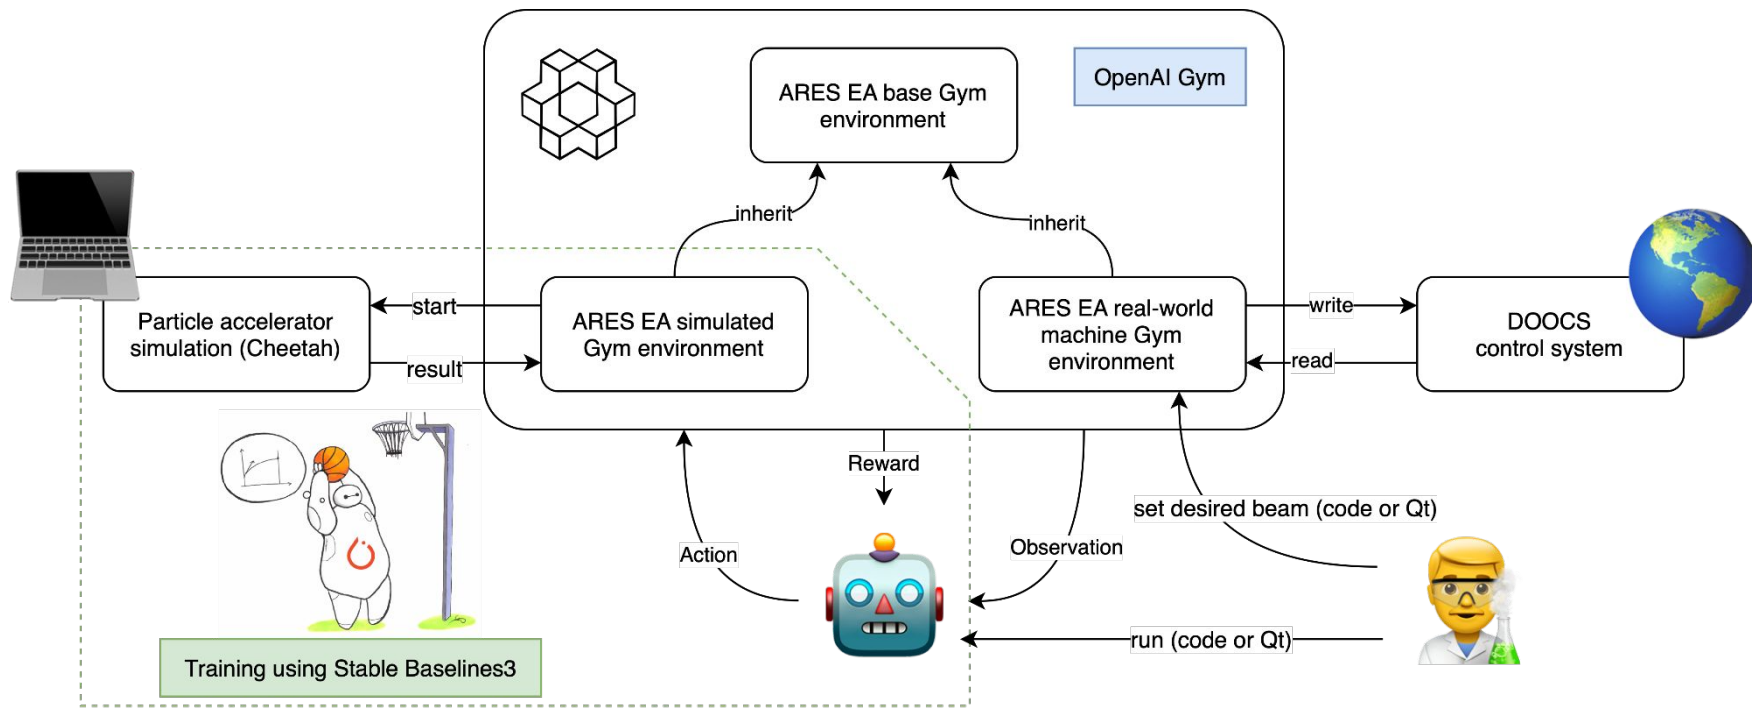
# Accelerator Example

## Positioning and focusing in the ARES Experimental Area





# Accelerator Example

## Technical overview




# Accelerator Example


## Training results



Q1=10.00 Q2=-10.00 CV=0.00 Q3=10.00  
CH=0.00  
mx=-0.07 sx=0.24 my=1.65 sy=0.66




Q1=6.31 Q2=-16.51 CV=0.00 Q3=16.98  
CH=0.00  
mx=0.03 sx=0.02 my=-0.22 sy=0.06



Q1=4.59 Q2=-14.39 CV=0.00 Q3=19.12  
CH=0.00  
mx=0.10 sx=0.02 my=-0.10 sy=0.03

# Accelerator Example

Running on the real accelerator



# Accelerator Example

## Getting it to work

### Choosing rewards

- Make beam as small as possible (when reading screen)  
➡ **Squeeze beam into corner**
- Minimise sum of pixels  
➡ **Push beam off-screen**
- Get positive reward for each beam parameter while in threshold. After 5 steps in threshold give win bonus and stop.
  - **Briefly jump out of threshold after 4 steps**



$$R(s_t, \mathbf{a}_t) = \begin{cases} \hat{R}(s_t, \mathbf{a}_t) & \text{if } \hat{R}(s_t, \mathbf{a}_t) > 0 \\ \underline{2 \cdot \hat{R}(s_t, \mathbf{a}_t)} & \text{otherwise.} \end{cases}$$

$$\hat{R}(s_t, \mathbf{a}_t) = O(\mathbf{x}_t) - O(\mathbf{x}_{t+1})$$

$$O(\mathbf{x}_t) = \underline{\ln} \sum_{p \in \mathbf{b}_t, p' \in \mathbf{b}'} w_p |p - p'|$$

### Sim2real transfer

Getting RL to run on simulations is easy. Getting it to run on a real accelerator is hard.

#### Domain randomisation




The simulation is never quite like the real world.

-> Add random noise.

#### Delta actions

Magnet settings may be affected by noise.

-> Policy outputs changes to magnet settings.



# Further Resources

## Where to start if you want to get into reinforcement learning

### Getting started in RL

- [OpenAI Spinning Up](#) - Very understandable explanations on RL and the most popular algorithms accompanied by easy-to-read Python implementations.
- [Reinforcement Learning with Stable Baselines 3](#) - YouTube playlist giving a good introduction on RL using Stable Baselines3.
- [Build a Doom AI Model with Python](#) - Detailed 3h tutorial of applying RL using *DOOM* as an example.
- [An introduction to Reinforcement Learning](#) - Brief introduction to RL.
- [An introduction to Policy Gradient methods](#) - Deep Reinforcement Learning - Brief introduction to PPO.
- [Real-time artificial intelligence for accelerator control: A study at the Fermilab Booster](#) - Regulation of a gradient magnet power supply using RL and real-time implementation of the trained agent using field-programmable gate arrays (FPGAs).
- [Magnetic control of tokamak plasmas through deep reinforcement learning](#) - Landmark paper on RL for controlling a real-world physical system (plasma in a tokamak fusion reactor).

### Papers

- [Learning-based optimisation of particle accelerators under partial observability without real-world training](#) - Tuning of electron beam properties on a diagnostic screen using RL.
- [Sample-efficient reinforcement learning for CERN accelerator control](#) - Beam trajectory steering using RL with a focus on sample-efficient training.
- [Autonomous control of a particle accelerator using deep reinforcement learning](#) - Beam transport through a drift tube linac using RL.
- [Basic reinforcement learning techniques to control the intensity of a seeded free-electron laser](#) - RL-based laser alignment and drift recovery.

### Literature

- [Reinforcement Learning: An Introduction](#) - Standard text book on RL.

### Packages

- [Gym](#) - Defacto standard for implementing custom environments. Also provides a library of RL tasks widely used for benchmarking.
- [Stable Baselines3](#) - Provides reliable, benchmarked and easy-to-use implementations of the most important RL algorithms.
- [Ray RLlib](#) - Part of the *Ray* Python package providing implementations of various RL algorithms with a focus on distributed training.

# Questions or remarks?



## Contact

**DESY.** Deutsches  
Elektronen-Synchrotron

[www.desy.de](http://www.desy.de)

Jan Kaiser & Oliver Stein  
Machine Beam Control (MSK)  
[jan.kaiser@desy.de](mailto:jan.kaiser@desy.de) | [oliver.stein@desy.de](mailto:oliver.stein@desy.de)